

Network System Design

CS6100

Tutorial 03

Hardware-based Multicast Filtering

Student Details

Name: R Abinav
Roll No: ME23B1004

Date: January 31, 2026

1 Problem Statement

In high-speed networks, a host interface receives significant multicast traffic. Processing every packet in the Operating System kernel to determine relevance is computationally expensive and causes high CPU interrupt overhead.

Network Interface Cards (NICs) employ a hardware-based "pre-filter" using hashing to discard unwanted packets before CPU interruption. This assignment simulates the hardware filtering mechanism and the complete life-cycle of a multicast packet from physical wire to final software decision.

1.1 Objectives

- **Understand CRC-32:** Implement CRC-32 hashing logic used in Ethernet standards
- **Simulate Hardware Constraints:** Model limited-size hardware hash table (16-64 entries)
- **Analyze Collisions:** Demonstrate false positives where unwanted traffic leaks through
- **Performance Metrics:** Calculate filtering ratio (hardware vs software processing)

2 Background Theory

2.1 Multicast Traffic and CPU Overhead

Network interfaces receive all packets on the wire, including multicast traffic for other hosts. Without hardware filtering, every packet triggers CPU interrupt for software processing, creating overhead from context switches and header examination. High multicast traffic can saturate CPU with interrupt handling.

2.2 Two-Stage Filtering Architecture

Modern NICs use two-stage filtering:

1. **Hardware Pre-filter:** Fast hash table on NIC drops most unwanted packets before CPU interruption
2. **Software Filter:** OS kernel maintains exact list of subscribed multicast groups for final verification

This trades some false positives for massive reduction in CPU interrupts.

2.3 CRC-32 Hash Function

CRC-32 (Cyclic Redundancy Check) is a polynomial-based hash function standardized in IEEE 802.3:

- Deterministic: same input produces same hash
- Fast: uses table lookup

- Good distribution: spreads values across hash space
- Hardware-friendly: simple bitwise operations
- Polynomial: 0xEDB88320 (reversed representation)

2.4 IP Multicast to MAC Mapping

IPv4 multicast (224.0.0.0 to 239.255.255.255) maps to Ethernet MAC per RFC 1112:

- Multicast MAC prefix: 01:00:5E
- Lower 23 bits of IP map to lower 23 bits of MAC
- Upper 5 bits lost in mapping
- Creates inherent collisions: 32 IPs map to each MAC

Example: 224.0.0.1 → 01:00:5E:00:00:01

2.5 Hash Collisions and False Positives

Hardware hash table is small (16-64 entries) to fit in fast NIC memory, creating:

1. **IP-to-MAC collisions:** Multiple IPs map to same MAC (RFC 1112 inherent)
2. **Hash collisions:** Multiple MACs hash to same table index

Collisions cause hardware filter to pass unwanted packets (false positives). Software filter performs exact matching to drop these.

3 Implementation

3.1 CRC-32 Hash Engine

- Precomputes 256-entry lookup table using polynomial 0xEDB88320
- Computes 32-bit CRC for 6-byte MAC addresses
- Extracts hash index using upper N bits (N = table size in bits)

3.2 Hardware Hash Table

- Bit vector (each bit = one hash bucket)
- Size: 4 bits = 16 entries
- Operations: set bit on subscription, check bit on packet arrival
- Passes packet if bit set (includes false positives)

3.3 Software Filter

- Maintains exact set of subscribed IP addresses
- Final verification on packets passed by hardware
- Drops false positives (unsubscribed IPs that passed hardware)

3.4 Packet Flow

1. Packet arrives with destination IP
2. Convert IP to multicast MAC
3. Hash MAC to get table index
4. Check hardware hash table bit
5. If bit = 0: drop (no CPU interrupt)
6. If bit = 1: pass to software
7. Software checks exact IP subscription
8. Accept if subscribed, drop if false positive

4 Simulation Results

Configuration:

- Hash table: 4 bits (16 entries)
- Subscribed groups: 5 multicast IPs
- Test packets: 400 (100 subscribed, 300 unsubscribed)

4.1 Output Analysis

4.2 Key Observations

1. Hardware Filtering Ratio: 60.00%

240/400 packets dropped by hardware without CPU interruption, representing 60% reduction in interrupt load.

2. False Positive Rate: 37.50%

60/160 packets passed hardware filter as false positives (unsubscribed traffic). These caused unnecessary CPU interrupts due to hash collisions.

3. Hash Collision Example

Well-known multicast addresses show multiple IPs mapping to same hash index:

- 224.0.0.1, 224.0.0.5, 224.0.0.9 → index 2

```

└─ $ ~/Desktop/course-work/nsd/tutorial-03 └─ main !2 ?1 .....
  cargo run --release
  Compiling tutorial-03 v0.1.0 (/Users/abinav/Desktop/course-work/nsd/tutorial-03)
    Finished `release` profile [optimized] target(s) in 1.15s
      Running `target/release/tutorial-03`
multicast filter simulation

hardware hash table size: 4 bits (16 entries)

subscribed multicast groups:
  224.0.0.1 -> 01:00:5E:00:00:01
  224.0.0.5 -> 01:00:5E:00:00:05
  224.0.0.251 -> 01:00:5E:00:00:FB
  239.192.1.1 -> 01:00:5E:40:01:01
  239.192.2.2 -> 01:00:5E:40:02:02

processing 400 packets...

simulation results:
  total packets: 400
  hardware dropped: 240
  hardware passed: 160
  software accepted: 100
  software dropped: 60

performance metrics:
  hardware filtering ratio: 60.00%
  false positive rate: 37.50%

hash collision analysis:

false positive examples (unsubscribed traffic leaked through hw filter):
  60 packets from unsubscribed groups passed hardware filter
  these share hash indices with subscribed groups

well-known multicast addresses:
  224.0.0.1 -> 01:00:5E:00:00:01 (hash index: 2)
  224.0.0.2 -> 01:00:5E:00:00:02 (hash index: 11)
  224.0.0.5 -> 01:00:5E:00:00:05 (hash index: 2)
  224.0.0.6 -> 01:00:5E:00:00:06 (hash index: 11)
  224.0.0.9 -> 01:00:5E:00:00:09 (hash index: 2)

```

Figure 1: Simulation output showing filtering statistics and hash collisions

- 224.0.0.2, 224.0.0.6 → index 11
4. **Software Accuracy: 100%**
 All 100 subscribed packets correctly accepted after passing hardware filtering.

5 Trade-offs and Design Considerations

5.1 Hash Table Size

- **Smaller (4 bits):** Less hardware cost, higher collisions, more false positives
- **Larger (6-8 bits):** Lower collisions, requires more NIC memory

5.2 Performance Impact

Despite 37.5% false positive rate, the system achieves:

- 60% reduction in CPU interrupts (240 packets dropped in hardware)
- Only 60 unnecessary interrupts (false positives) vs 300 without filtering
- Net benefit: 80% reduction in unnecessary CPU load (240 vs 300)

5.3 Real-world Implementation

Production NICs typically use:

- 64-entry hash tables (6 bits) for cost-performance balance
- Perfect hash functions optimized for common multicast patterns
- Additional filtering layers (VLAN tags, port filtering)

6 Conclusion

The simulation demonstrates hardware-based multicast filtering reduces CPU overhead despite hash collisions:

1. CRC-32 provides fast, deterministic hashing for NIC hardware
2. Small hash tables create collisions but provide significant benefit
3. Two-stage architecture balances hardware cost with CPU efficiency
4. False positives are acceptable for dramatic interrupt reduction

This filtering mechanism is essential for high-speed networking, enabling hosts to handle multicast-heavy traffic without CPU saturation.

7 Source Code

```

1  use std::collections::{HashMap, HashSet};
2  use std::fmt;
3
4  //crc32 ieee 802.3
5  struct Crc32 {
6      table: [u32; 256],
7  }
8
9  impl Crc32 {
10     const POLYNOMIAL: u32 = 0xEDB88320;
11
12     fn new() -> Self {
13         let mut table = [0u32; 256];
14         for i in 0..256 {
15             let mut crc = i as u32;
16             for _ in 0..8 {
17                 if crc & 1 != 0 {
18                     crc = (crc >> 1) ^ Self::POLYNOMIAL;
19                 } else {
20                     crc >>= 1;
21                 }
22             }
23             table[i] = crc;
24         }
25         Self { table }
26     }
27
28     fn compute(&self, data: &[u8]) -> u32 {
29         let mut crc = 0xFFFFFFFF;
30         for &byte in data {
31             let index = ((crc ^ byte as u32) & 0xFF) as usize;
32             crc = (crc >> 8) ^ self.table[index];
33         }
34         !crc
35     }
36
37     fn hash_to_index(&self, mac: &MacAddress, bits: u8) -> usize {
38         let crc = self.compute(&mac.0);
39         let shift = 32 - bits;
40         ((crc >> shift) as usize) & ((1 << bits) - 1)
41     }
42 }
43
44 #[derive(Debug, Clone, Copy, PartialEq, Eq, Hash)]
45 struct MacAddress([u8; 6]);
46
47 impl MacAddress {
48     fn from_multicast_ip(ip: &IpAddress) -> Self {
49         let b = ip.0;
50         Self([0x01, 0x00, 0x5E, b[1] & 0x7F, b[2], b[3]])
51     }
52 }
53
54 impl fmt::Display for MacAddress {
55     fn fmt(&self, f: &mut fmt::Formatter) -> fmt::Result {

```

```

56         write!(  
57             f,  
58             "{:02X}:{:02X}:{:02X}:{:02X}:{:02X}:{:02X}",  
59             self.0[0], self.0[1], self.0[2], self.0[3], self.0[4],  
60             self.0[5]  
61         )  
62     }  
63  
64 #[derive(Debug, Clone, Copy, PartialEq, Eq, Hash)]  
65 struct IpAddress([u8; 4]);  
66  
67 impl IpAddress {  
68     fn new(a: u8, b: u8, c: u8, d: u8) -> Self {  
69         Self([a, b, c, d])  
70     }  
71  
72     fn is_multicast(&self) -> bool {  
73         self.0[0] >= 224 && self.0[0] <= 239  
74     }  
75 }  
76  
77 impl fmt::Display for IpAddress {  
78     fn fmt(&self, f: &mut fmt::Formatter) -> fmt::Result {  
79         write!(f, "{}.{}.{}.{}", self.0[0], self.0[1], self.0[2],  
80             self.0[3])  
81     }  
82 }  
83 #[derive(Clone)]  
84 struct MulticastPacket {  
85     dst_ip: IpAddress,  
86     dst_mac: MacAddress,  
87 }  
88  
89 impl MulticastPacket {  
90     fn new(dst_ip: IpAddress) -> Self {  
91         let dst_mac = MacAddress::from_multicast_ip(&dst_ip);  
92         Self { dst_ip, dst_mac }  
93     }  
94 }  
95  
96 struct HardwareHashTable {  
97     bits: Vec<bool>,  
98     size_bits: u8,  
99     crc: Crc32,  
100 }  
101  
102 impl HardwareHashTable {  
103     fn new(size_bits: u8) -> Self {  
104         Self {  
105             bits: vec![false; 1 << size_bits],  
106             size_bits,  
107             crc: Crc32::new(),  
108         }  
109     }  
110  
111     fn add_mac(&mut self, mac: &MacAddress) {
```

```

112         let idx = self.crc.hash_to_index(mac, self.size_bits);
113         self.bits[idx] = true;
114     }
115
116     fn check_mac(&self, mac: &MacAddress) -> bool {
117         let idx = self.crc.hash_to_index(mac, self.size_bits);
118         self.bits[idx]
119     }
120 }
121
122 struct SoftwareFilter {
123     subscribed: HashSet<IpAddress>,
124 }
125
126 impl SoftwareFilter {
127     fn new() -> Self {
128         Self {
129             subscribed: HashSet::new(),
130         }
131     }
132
133     fn subscribe(&mut self, ip: IpAddress) {
134         self.subscribed.insert(ip);
135     }
136
137     fn is_subscribed(&self, ip: &IpAddress) -> bool {
138         self.subscribed.contains(ip)
139     }
140 }
141
142 #[derive(Default)]
143 struct SimulationStats {
144     total: usize,
145     hw_dropped: usize,
146     hw_passed: usize,
147     sw_accepted: usize,
148     sw_dropped: usize,
149 }
150
151 struct MulticastFilterSimulator {
152     hw: HardwareHashTable,
153     sw: SoftwareFilter,
154     stats: SimulationStats,
155     mac_to_ips: HashMap<MacAddress, Vec<IpAddress>>,
156 }
157
158 impl MulticastFilterSimulator {
159     fn new(bits: u8) -> Self {
160         Self {
161             hw: HardwareHashTable::new(bits),
162             sw: SoftwareFilter::new(),
163             stats: SimulationStats::default(),
164             mac_to_ips: HashMap::new(),
165         }
166     }
167
168     fn subscribe(&mut self, ip: IpAddress) {
169         assert!(ip.is_multicast());

```

```

170         self.sw.subscribe(ip);
171         let mac = MacAddress::from_multicast_ip(&ip);
172         self.hw.add_mac(&mac);
173         self.mac_to_ips.entry(mac).or_insert_with(Vec::new).push(ip)
174             ;
175     }
176
177     fn process(&mut self, pkt: MulticastPacket) {
178         self.stats.total += 1;
179
180         if !self.hw.check_mac(&pkt.dst_mac) {
181             self.stats.hw_dropped += 1;
182             return;
183         }
184
185         self.stats.hw_passed += 1;
186
187         if self.sw.is_subscribed(&pkt.dst_ip) {
188             self.stats.sw_accepted += 1;
189         } else {
190             self.stats.sw_dropped += 1;
191         }
192     }
193
194     fn generate_well_known_addresses() -> Vec<IpAddress> {
195         vec![
196             IpAddress::new(224, 0, 0, 1),
197             IpAddress::new(224, 0, 0, 2),
198             IpAddress::new(224, 0, 0, 5),
199             IpAddress::new(224, 0, 0, 6),
200             IpAddress::new(224, 0, 0, 9),
201         ]
202     }
203
204     fn main() {
205         let mut sim = MulticastFilterSimulator::new(4);
206
207         let subs = vec![
208             IpAddress::new(224, 0, 0, 1),
209             IpAddress::new(224, 0, 0, 5),
210             IpAddress::new(224, 0, 0, 251),
211             IpAddress::new(239, 192, 1, 1),
212             IpAddress::new(239, 192, 2, 2),
213         ];
214
215         println!("multicast filter simulation");
216         println!();
217         println!("hardware hash table size: {} bits ({} entries)", 4, 1
218             << 4);
219         println!();
220
221         println!("subscribed multicast groups:");
222         for ip in &subs {
223             let mac = MacAddress::from_multicast_ip(ip);
224             println!("  {} -> {}", ip, mac);
225             sim.subscribe(*ip);
226         }

```

```

226     println!();
227
228     let mut packets = Vec::new();
229
230     for ip in &subs {
231         for _ in 0..20 {
232             packets.push(MulticastPacket::new(*ip));
233         }
234     }
235
236     let others = vec![
237         IpAddress::new(224, 0, 0, 2),
238         IpAddress::new(224, 0, 0, 9),
239         IpAddress::new(224, 0, 1, 1),
240         IpAddress::new(239, 100, 1, 1),
241         IpAddress::new(239, 100, 2, 2),
242         IpAddress::new(238, 50, 50, 50),
243         IpAddress::new(224, 1, 1, 1),
244         IpAddress::new(224, 2, 2, 2),
245         IpAddress::new(225, 1, 1, 1),
246         IpAddress::new(230, 5, 5, 5),
247     ];
248
249     for ip in &others {
250         for _ in 0..30 {
251             packets.push(MulticastPacket::new(*ip));
252         }
253     }
254
255     println!("processing {} packets...", packets.len());
256     println!();
257
258     for pkt in packets {
259         sim.process(pkt);
260     }
261
262     println!("simulation results:");
263     println!("    total packets: {}", sim.stats.total);
264     println!("    hardware dropped: {}", sim.stats.hw_dropped);
265     println!("    hardware passed: {}", sim.stats.hw_passed);
266     println!("    software accepted: {}", sim.stats.sw_accepted);
267     println!("    software dropped: {}", sim.stats.sw_dropped);
268     println!();
269
270     let hw_filter_rate = (sim.stats.hw_dropped as f64 / sim.stats.
271                           total as f64) * 100.0;
272     let false_positive_rate = (sim.stats.sw_dropped as f64 / sim.
273                               stats.hw_passed as f64) * 100.0;
274
275     println!("performance metrics:");
276     println!("    hardware filtering ratio: {:.2}%", hw_filter_rate);
277     println!("    false positive rate: {:.2}%", false_positive_rate);
278     println!();
279
280     println!("hash collision analysis:");
281     for (mac, ips) in &sim.mac_to_ips {
282         if ips.len() > 1 {
283             println!("    collision at mac {}", mac);

```

```
282         for ip in ips {
283             println!("    {}", ip);
284         }
285     }
286
287     if sim.stats.sw_dropped > 0 {
288         println!();
289         println!("false positive examples (unsubscribed traffic
290             leaked through hw filter):");
291         println!("    {} packets from unsubscribed groups passed
292             hardware filter", sim.stats.sw_dropped);
293         println!("    these share hash indices with subscribed groups
294             ");
295     }
296
297     let crc = Crc32::new();
298     println!();
299     println!("well-known multicast addresses:");
300     for ip in generate_well_known_addresses().iter() {
301         let mac = MacAddress::from_multicast_ip(ip);
302         let hash_index = crc.hash_to_index(&mac, 4);
303         println!("    {} -> {} (hash index: {})", ip, mac, hash_index)
304     }
305 }
```